

Accelerating Neural Networks on FPGAs

- Moiz Zaheer Malik
- B.Eng. Electronic Engineering
- Hochschule Hamm-Lippstadt



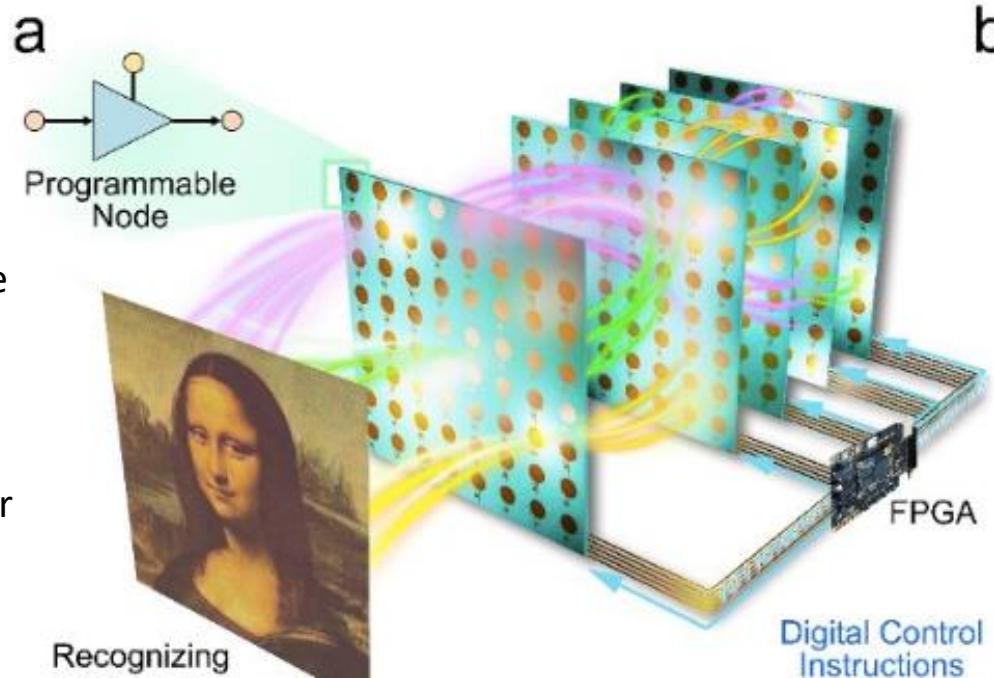
Agenda



Motivation

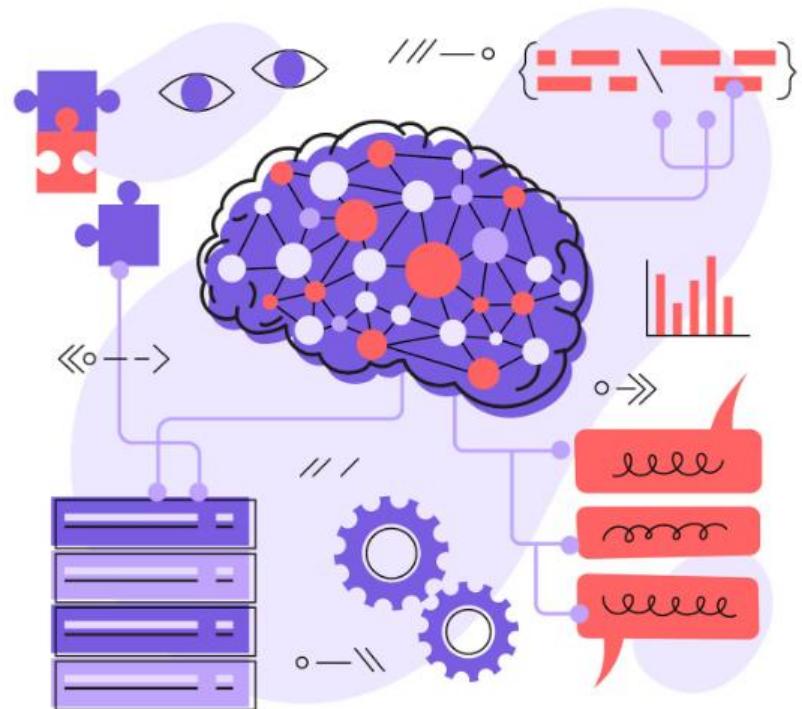
Neural networks are used everywhere

- Neural networks are used everywhere
- Image recognition
- Self driving cars
- Speech and language processing
- They need very high computing power



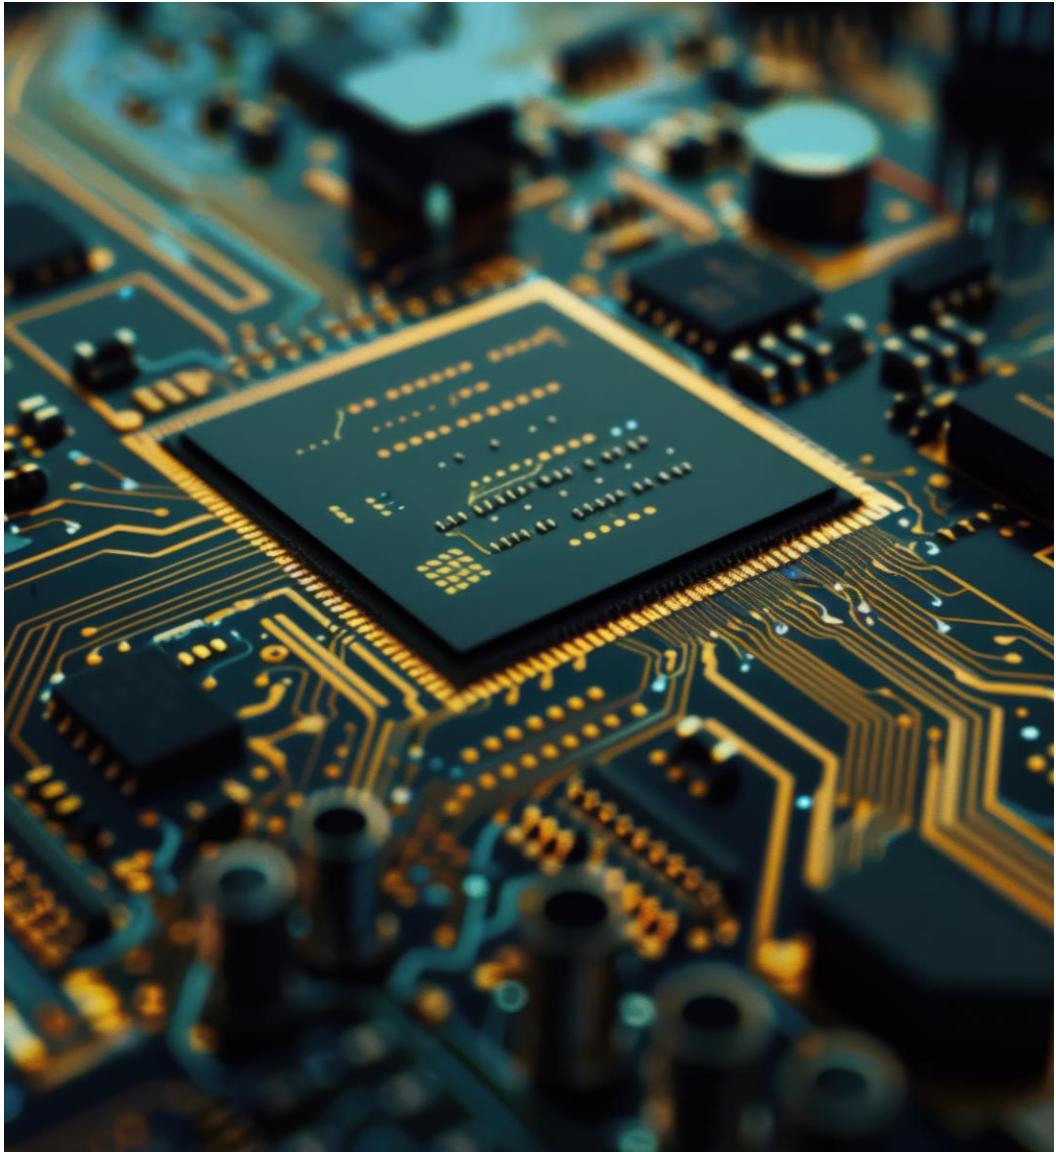
Problem Statement

- Modern neural networks are very large
- Millions of parameters
- Billions of operations
- CPUs are too slow
- GPUs consume a lot of power
- Not good for real time systems



What is an FPGA?

- FPGA is programmable hardware
- Custom logic design
- Parallel processing
- Low power consumption



```
# This program adds two numbers

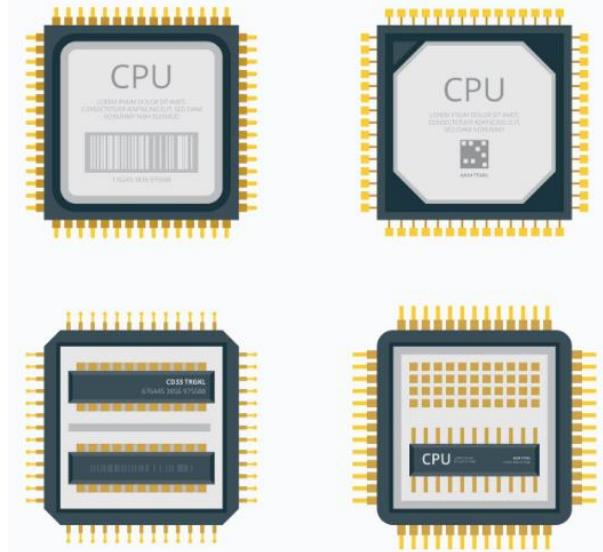
num1 = 1.5
num2 = 6.3

# Add two numbers
sum = num1 + num2

# Display the sum
print('The sum of {0} and {1} is {2}'.format(num1, num2, sum))
```



Compilation



- Code is written by the developer
- Code is compiled
- Compilation produces a set of instructions
- These instructions run on a processor
- The processor executes instructions one by one

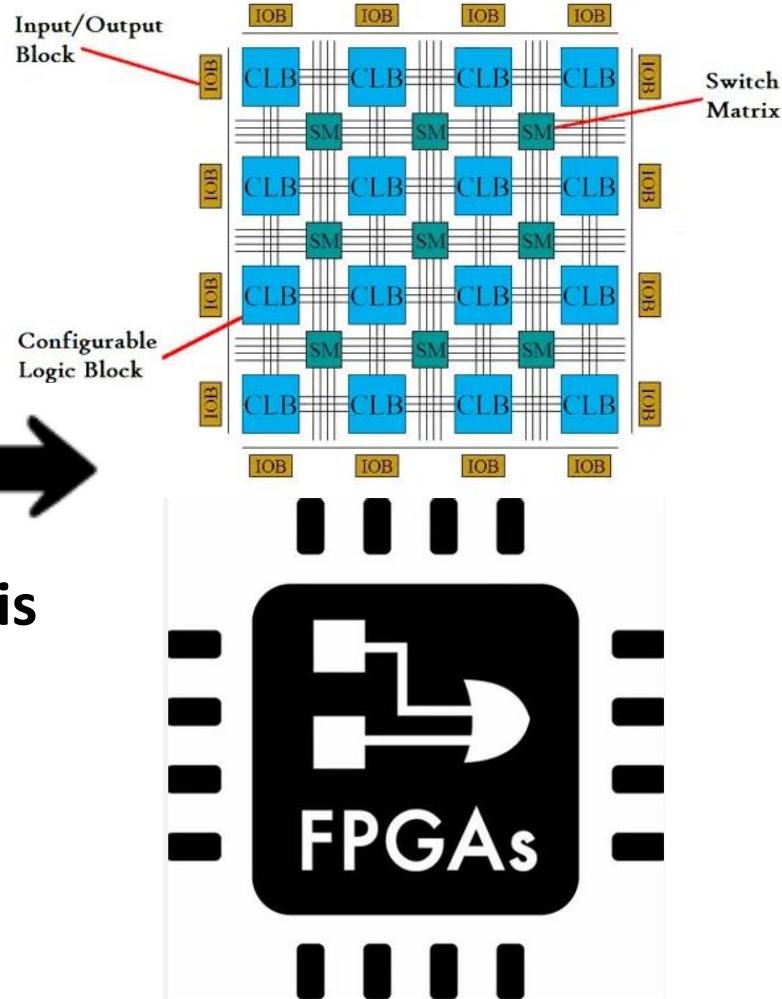
```

module adder (
    input wire [4:0] a,
    input wire [4:0] b,
    output wire [4:0] y
);
    assign y = a + b;
endmodule

```



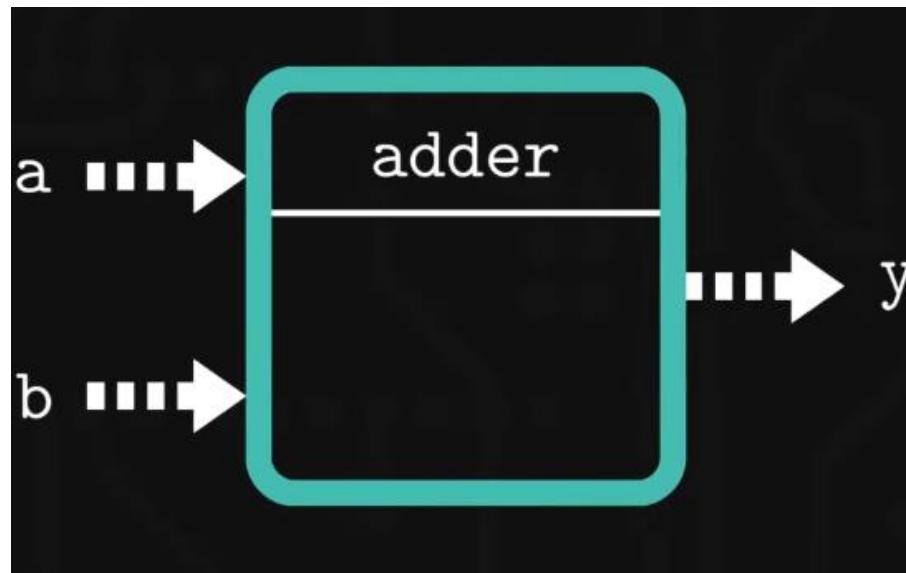
Synthesis



- Code can be written in a different way
- This code is synthesized for an FPGA
- The result is hardware, not instructions
- This type of code describes hardware behavior
- It is called Hardware Description Language (HDL)

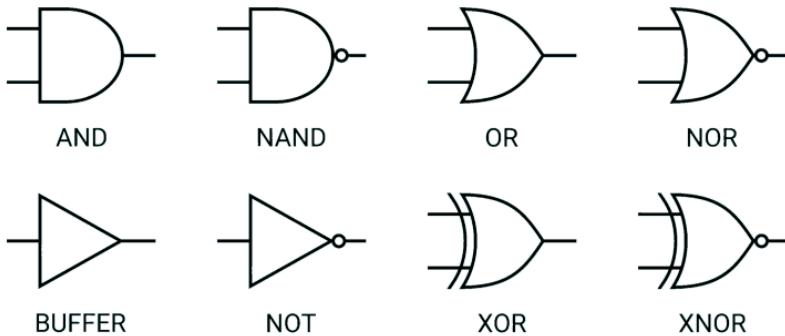
Synthesis

- Synthesis maps code to physical hardware
- Code describes operations, for example adding two numbers
- Inputs: A and B
- Output: Y
- Synthesizer creates hardware blocks for this logic
- The result is real hardware, not software instructions



Logic blocks inside an FPGA

Logic Gates

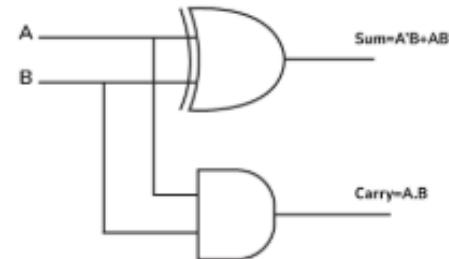


- Logic blocks are made of registers and logic gates
- Logic gates can be AND, OR, XNOR, NOT
- Gates are connected together to form logic blocks

Example logic blocks:

- Half adder
- Full adder

Half Adder

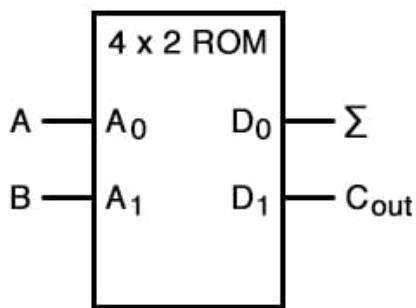


Truth Table

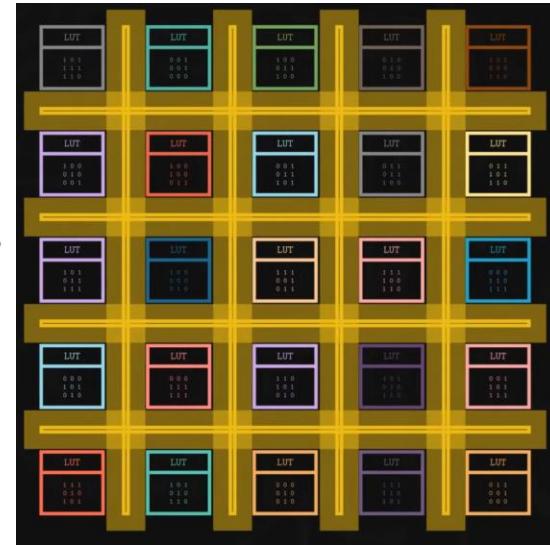
Inputs		Outputs	
A	B	Sum	Carry
0	0	0	0
0	1	1	0
1	0	1	0
1	1	0	1

Inputs go in, output comes out

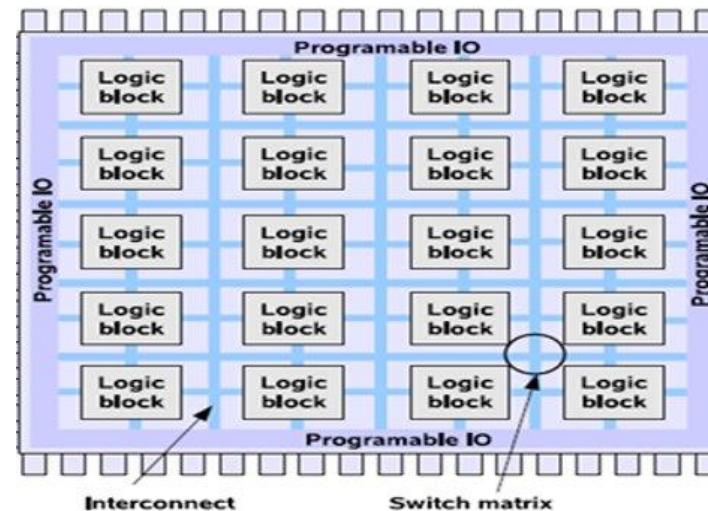
LUTs and FPGA structure



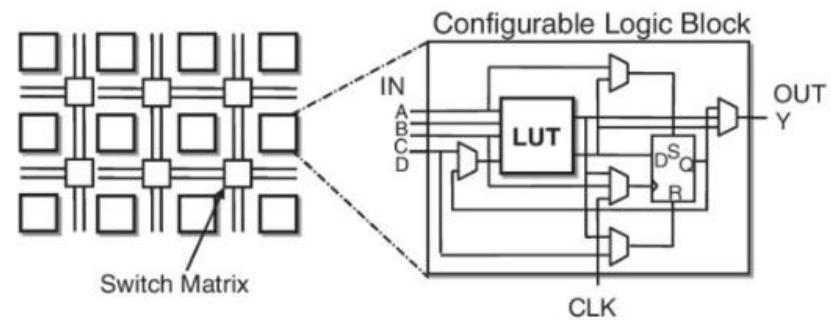
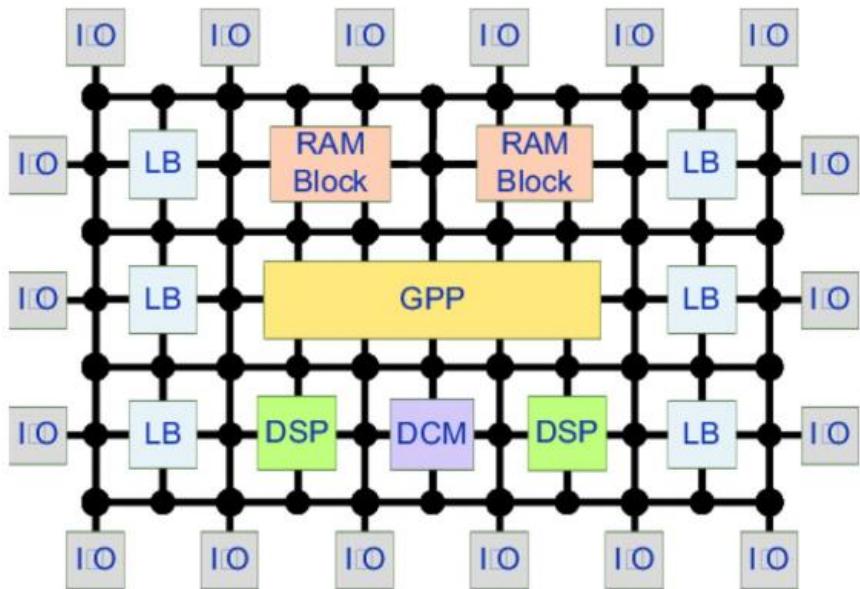
Address		Data	
A	B	C _{out}	Σ
0	0	0	0
0	1	0	1
1	0	0	1
1	1	1	0



- Logic gates are implemented using lookup tables (LUTs)
- A LUT implements logic using stored values
- Multiple LUTs are combined together
- LUTs are connected through a programmable switching fabric
- This interconnection is programmable
- All together, this forms an FPGA
- The FPGA behavior depends on how we code it

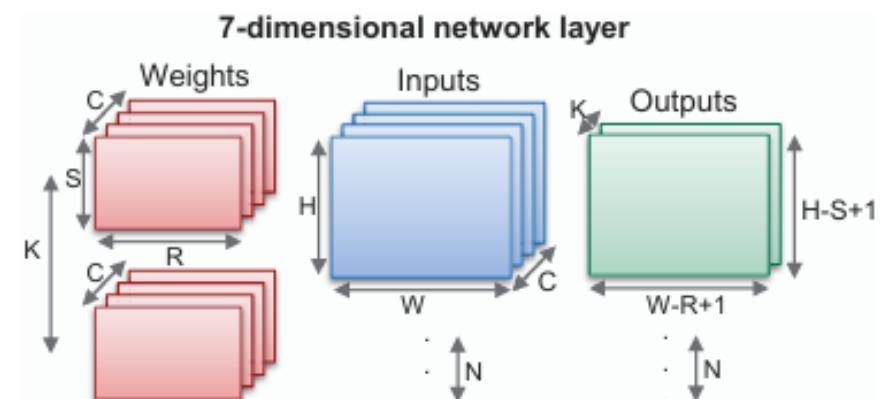
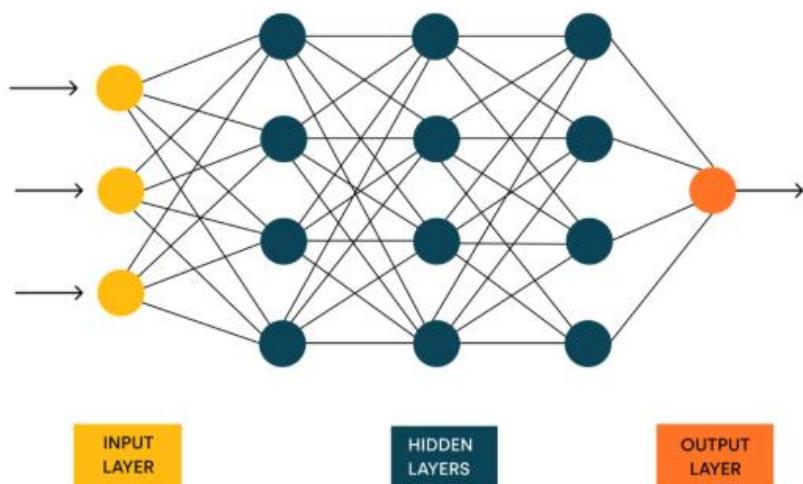


Inside an FPGA



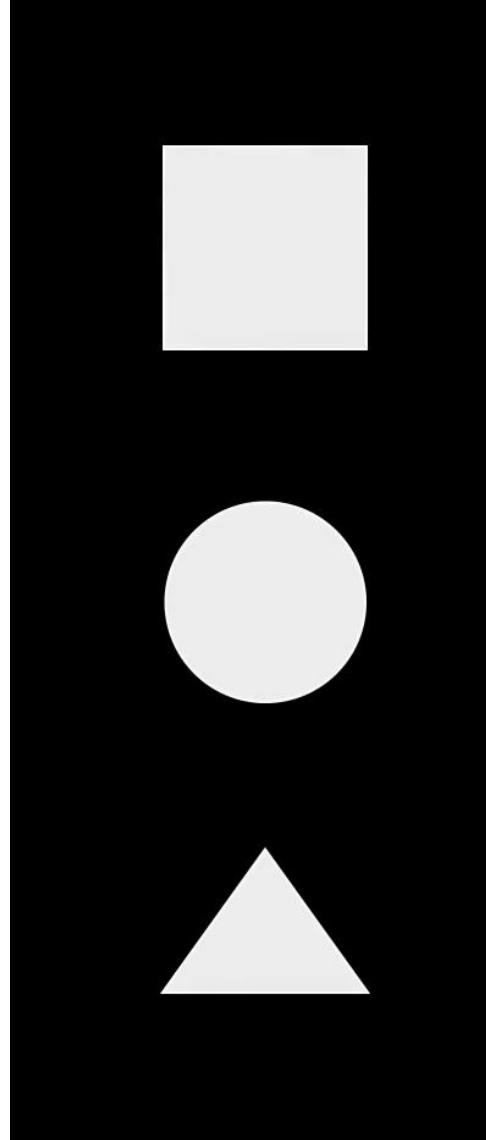
What is a Neural Network?

- Made of layers
- Uses weights
- Processes data step by step
- Produces output

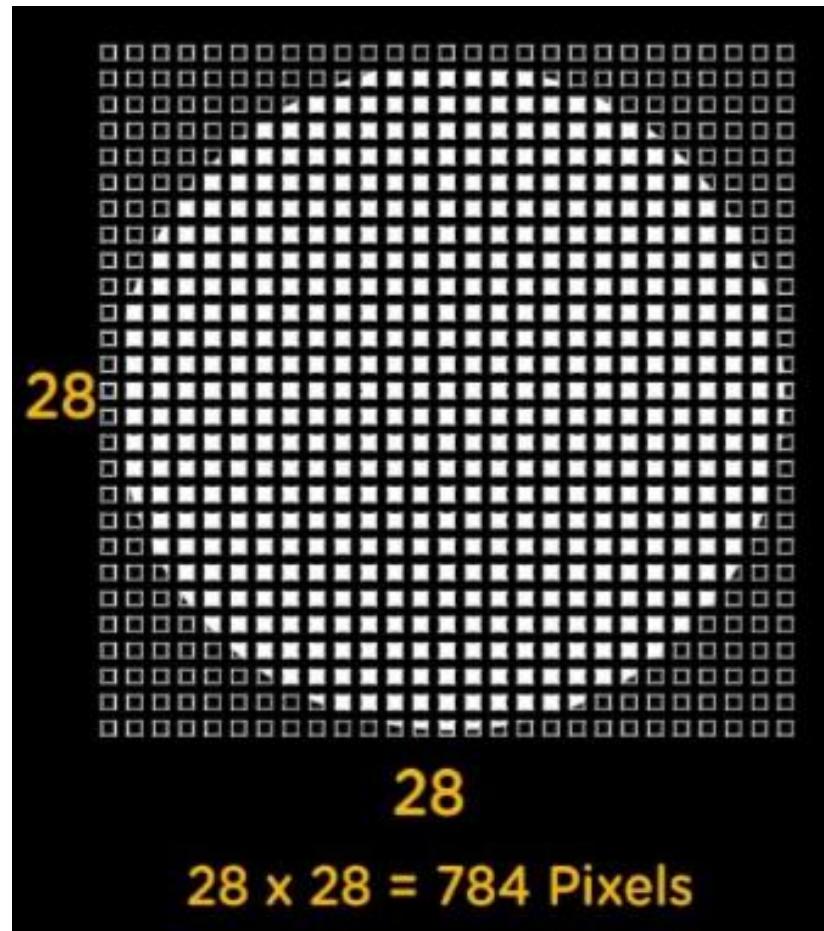




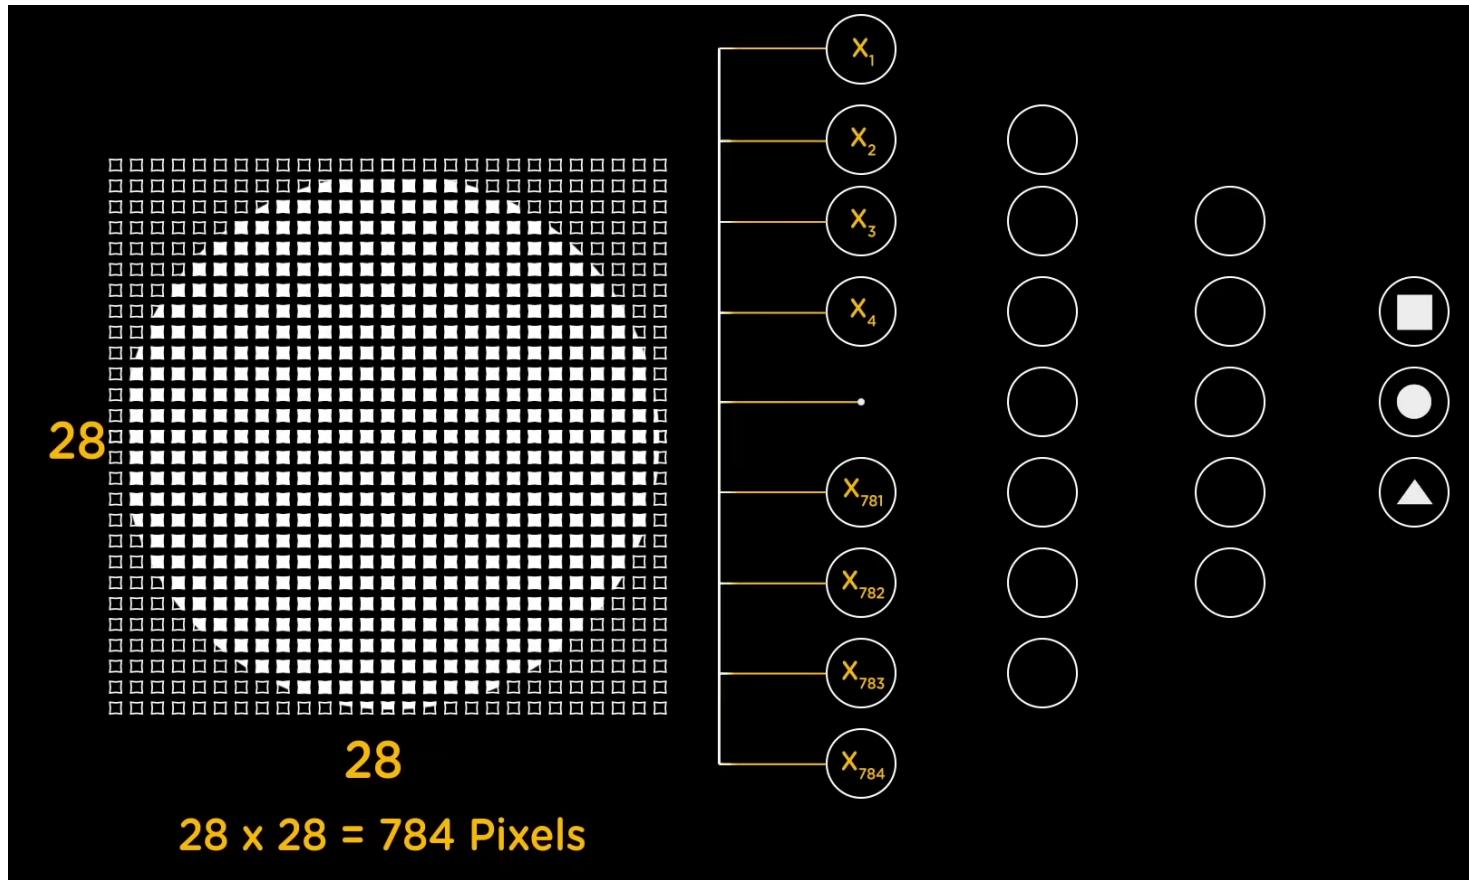
How Neural Networks Work



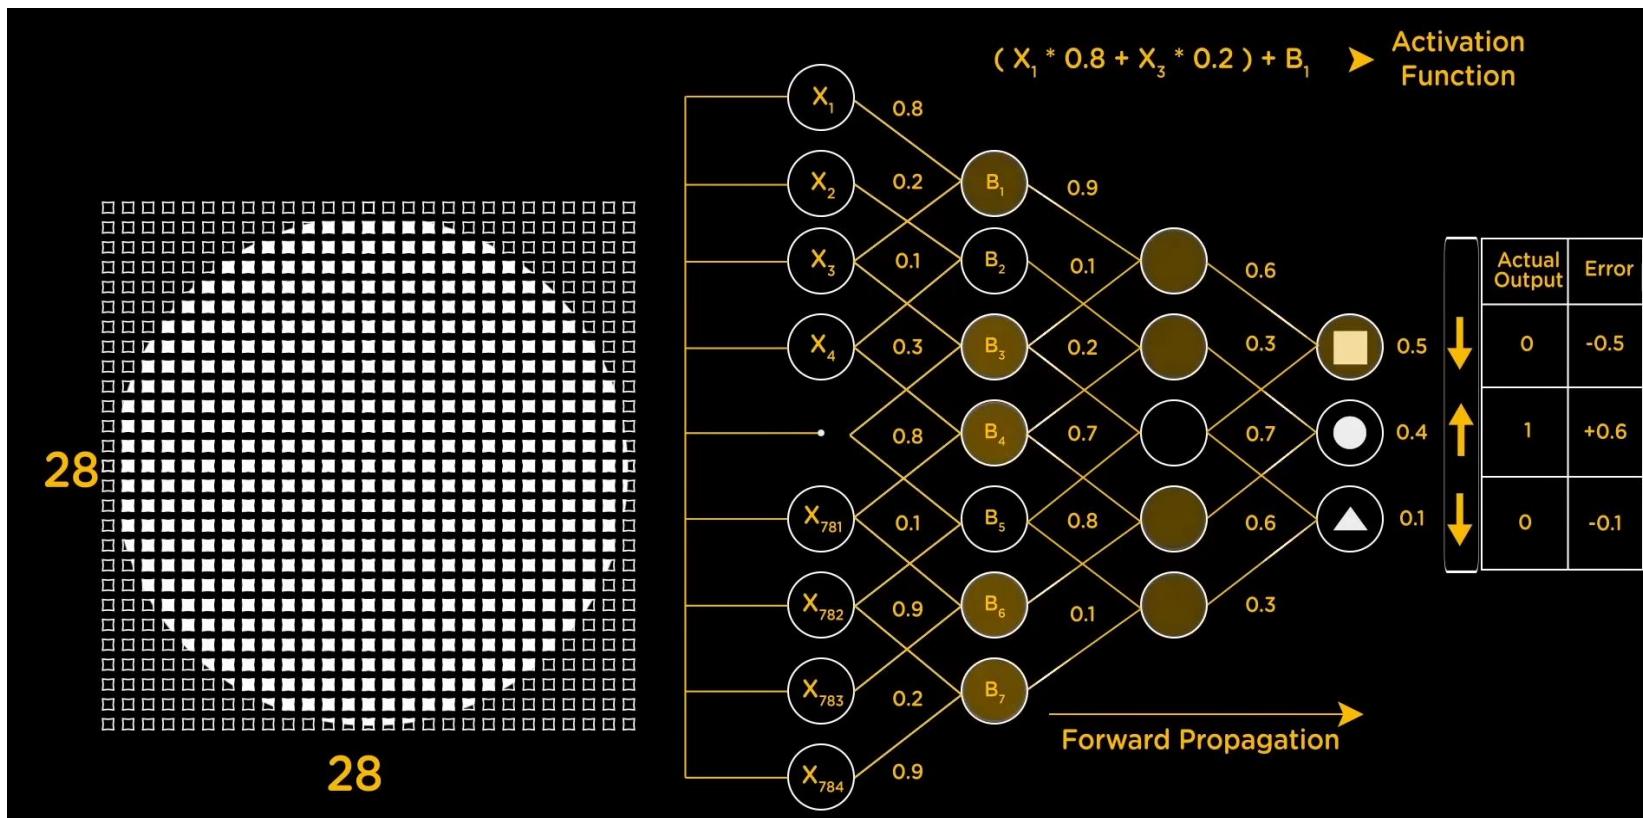
How Neural Networks Work



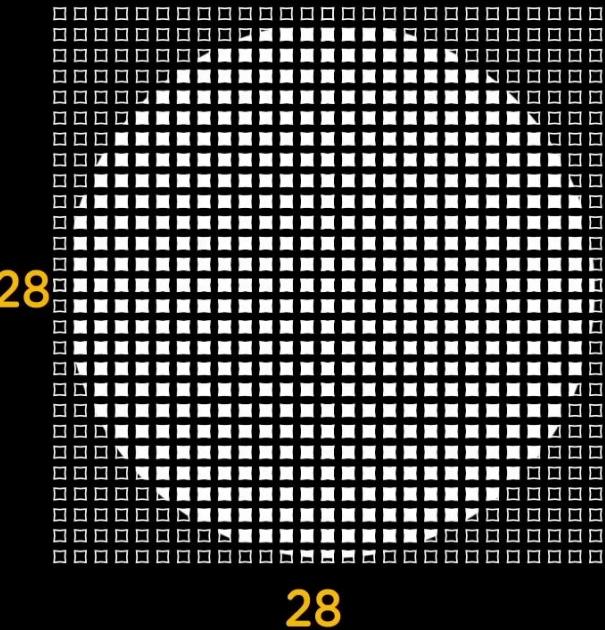
How Neural Networks Work



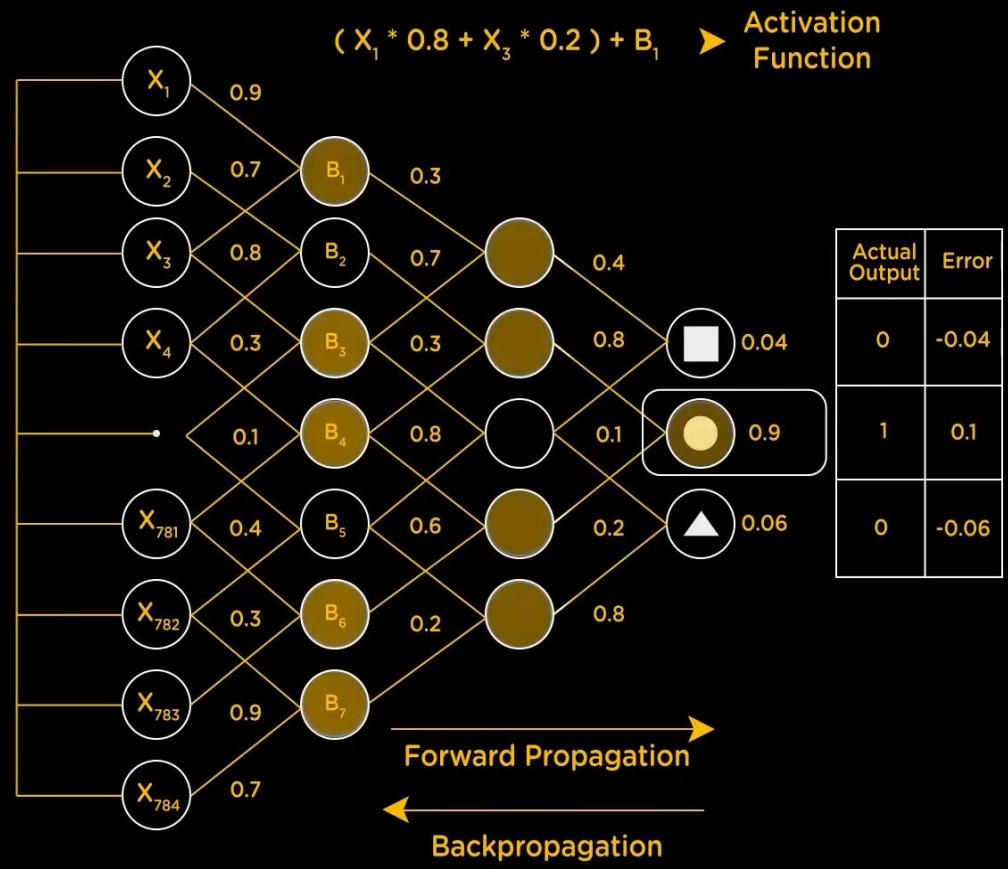
How Neural Networks Work



How Neural Networks Work



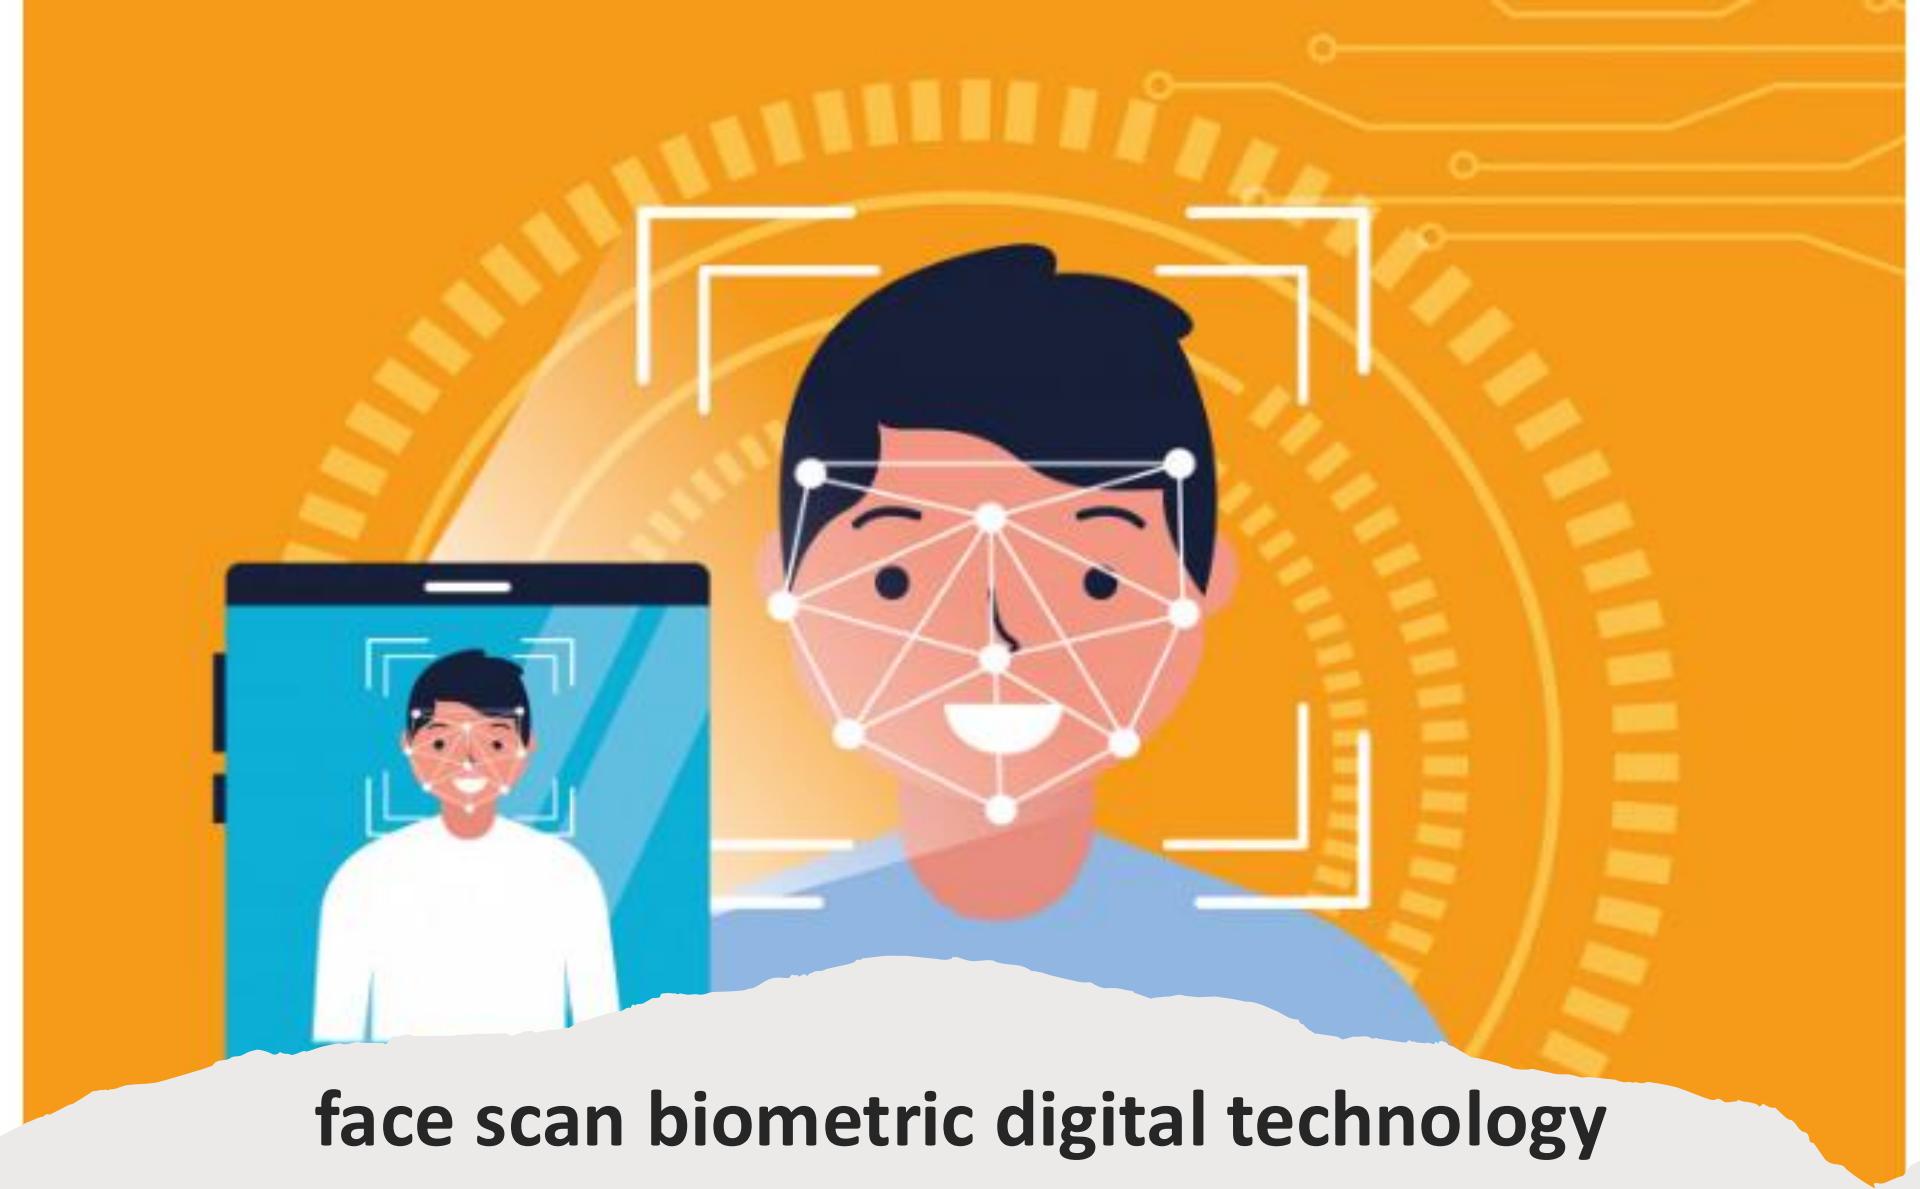
$28 \times 28 = 784$ Pixels





Question

Can anyone give an example of a neural network that they use in daily life from what we just learned?



face scan biometric digital technology

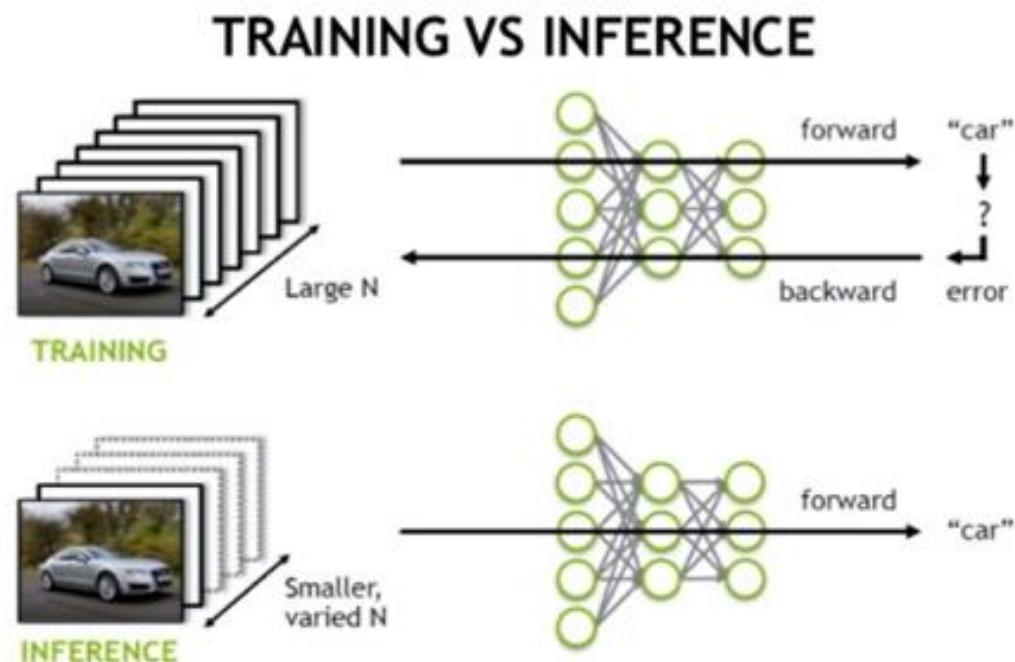
Training vs Inference

Training:

- Learn weights from data
- Very compute heavy

Inference:

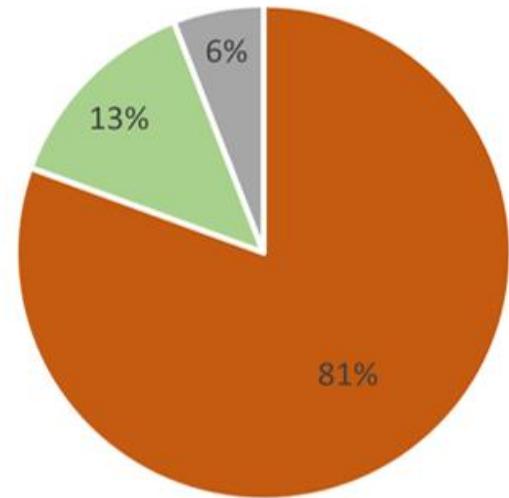
- Use trained model
- Only forward pass
- Used in real applications
- FPGA focus: inference



Why Use FPGA for Neural Networks?

- Parallel execution
- Low latency
- Low power consumption
- Reconfigurable hardware

- Model Inference
- Model training
- Matrix(Vector) Multiplication



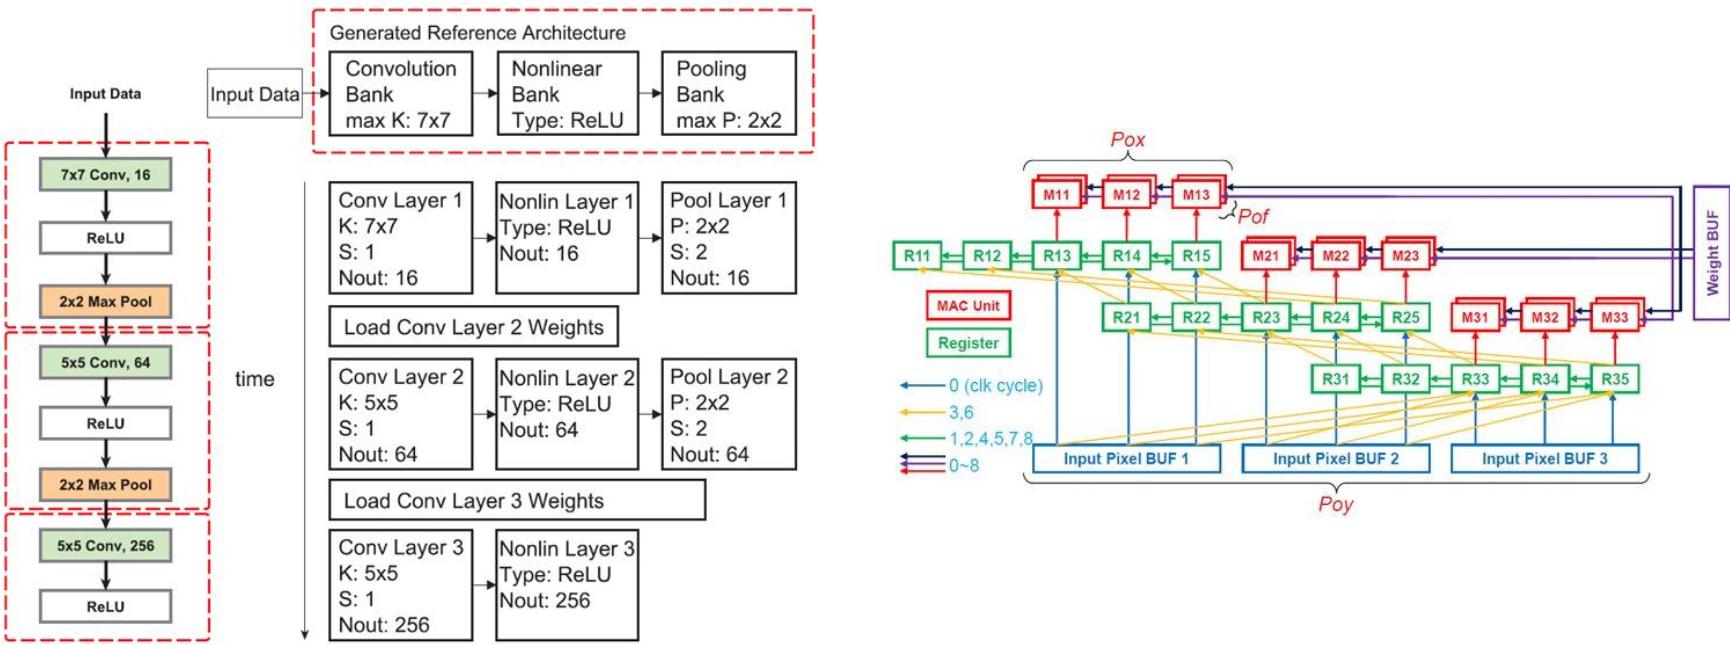
Main Acceleration Techniques

Parallelism

Pipelining

Quantization

Data reuse



Example FPGA Architectures

- Researchers build FPGA designs for neural networks
- Some designs are made for known networks like VGG16
- Some designs use math tricks to reduce work
- Data is reused to save time and power

Challenges



Limited hardware resources



Memory bandwidth limits



Design complexity



Precision trade-offs

Conclusion

FPGAs are good for neural networks

High efficiency and flexibility

Good for real-time and edge systems

Future work improves tools and memory



Any
Questions?

